

# THE SPARKS FOUNDATION: DATA SCIENCE AND BUSINESS ANALYTICS

Task 1: Prediction using Supervised ML

AIM: Predict the percentage of a student based on the number of study hours

AUTHOR: Nikita Nikam

LANGUAGE USED: Python 3

IDE: Jupyter Notebook

TYPE: Linear Regression

STEPS TO BE FOLLOWED:

Step 1: Import the Dataset

Step 2: Visualize and Analyze the Dataset

Step 3: Prepare the Data

Step 4: Design and Train the Machine Learning Model

Step 5: Visualize the Model

Step 6: Make Predictions

Step 7: Evaluate the Model

## Step 1: Import the Dataset

```
In [2]: import pandas as pd  
import numpy as np  
import matplotlib.pyplot as plt
```

In [3]:

```
# Reading data from remote link using the url

url = "http://bit.ly/w-data"
student_data = pd.read_csv(url)

print("Data imported successfully")
student_data
```

Data imported successfully

Out[3]:

	Hours	Scores
0	2.5	21
1	5.1	47
2	3.2	27
3	8.5	75
4	3.5	30
5	1.5	20
6	9.2	88
7	5.5	60
8	8.3	81
9	2.7	25
10	7.7	85
11	5.9	62
12	4.5	41
13	3.3	42
14	1.1	17
15	8.9	95
16	2.5	30
17	1.9	24
18	6.1	67
19	7.4	69
20	2.7	30
21	4.8	54
22	3.8	35
23	6.9	76
24	7.8	86

In [4]:

```
student_data.shape
```

*#Here we can see that there are 25 rows and 2 columns in the dataset*

Out[4]: (25, 2)

In [5]:

```
student_data.describe()
```

Out[5]:

	Hours	Scores
count	25.000000	25.000000
mean	5.012000	51.480000
std	2.525094	25.286887
min	1.100000	17.000000
25%	2.700000	30.000000
50%	4.800000	47.000000
75%	7.400000	75.000000
max	9.200000	95.000000

In [6]:

```
student_data.isnull().sum()
```

*#Here we can see that there are no NULL values in the dataset that can affect the*

Out[6]:


```
Hours      0
Scores     0
dtype: int64
```

## Step 2: Visualize and Analyze the Dataset

In [25]:

*# Plotting the distribution of scores and number of hours studied on a 2D graph*

```
student_data.plot(x='Hours', y='Scores', style='*',color='m')  
plt.title('Hours vs Percentage')  
plt.xlabel('Hours Studied')  
plt.ylabel('Percentage Score')  
plt.show()
```



From the above graph we can see that there is a positive linear relationship between hours and percentage which means that as the number of hours studied increased, the percentage scored also increased.

## Step 3: Prepare the Data

In [8]:

*# We are extracting values of Hours Data into variable X and the values of Scores*

```
X = student_data.iloc[:, :-1].values  
y = student_data.iloc[:, 1].values
```

```
In [9]: #Number of hours studied  
X
```

```
Out[9]: array([[2.5],  
               [5.1],  
               [3.2],  
               [8.5],  
               [3.5],  
               [1.5],  
               [9.2],  
               [5.5],  
               [8.3],  
               [2.7],  
               [7.7],  
               [5.9],  
               [4.5],  
               [3.3],  
               [1.1],  
               [8.9],  
               [2.5],  
               [1.9],  
               [6.1],  
               [7.4],  
               [2.7],  
               [4.8],  
               [3.8],  
               [6.9],  
               [7.8]])
```

```
In [10]: #Scores Obtained  
y
```

```
Out[10]: array([21, 47, 27, 75, 30, 20, 88, 60, 81, 25, 85, 62, 41, 42, 17, 95, 30,  
                24, 67, 69, 30, 54, 35, 76, 86], dtype=int64)
```

```
In [11]: # We now split the data into train and test datasets using Scikit-Learn's built-in  
  
from sklearn.model_selection import train_test_split  
  
X_train, X_test, y_train, y_test = train_test_split(X, y,  
                                                    test_size=0.2, random_state=0)
```

```
In [12]: X_train
```

```
Out[12]: array([[3.8],
                [1.9],
                [7.8],
                [6.9],
                [1.1],
                [5.1],
                [7.7],
                [3.3],
                [8.3],
                [9.2],
                [6.1],
                [3.5],
                [2.7],
                [5.5],
                [2.7],
                [8.5],
                [2.5],
                [4.8],
                [8.9],
                [4.5]])
```

```
In [13]: X_test
```

```
Out[13]: array([[1.5],
                [3.2],
                [7.4],
                [2.5],
                [5.9]])
```

```
In [14]: y_train
```

```
Out[14]: array([35, 24, 86, 76, 17, 47, 85, 42, 81, 88, 67, 30, 25, 60, 30, 75, 21,
                54, 95, 41], dtype=int64)
```

```
In [15]: y_test
```

```
Out[15]: array([20, 27, 69, 30, 62], dtype=int64)
```

## Step 4: Design and Train the Machine Learning Model

```
In [16]: from sklearn.linear_model import LinearRegression
```

```
regressor = LinearRegression()
regressor.fit(X_train, y_train)
```


```
print("Training complete.")
```

Training complete.

## Step 5: Visualize the Model

```
In [26]: # Plotting the regression line
line = regressor.coef_*X+regressor.intercept_

# Plotting for the test data
plt.scatter(X, y)
plt.plot(X, line, color='yellow');
plt.show()
```



## Step 6: Make Predictions

```
In [18]: print(X_test) # Testing data - In Hours
y_pred = regressor.predict(X_test) # Predicting the scores
```

```
[[1.5]
 [3.2]
 [7.4]
 [2.5]
 [5.9]]
```

```
In [19]: # Comparing Actual vs Predicted
df = pd.DataFrame({'Actual': y_test, 'Predicted': y_pred})
df
```

```
Out[19]:
```

	Actual	Predicted
0	20	16.884145
1	27	33.732261
2	69	75.357018
3	30	26.794801
4	62	60.491033

In [20]: *# Testing with custom data of 9.25 hrs/ day*

```
hours = 9.25
own_pred = regressor.predict([[hours]])
print(f"No of Hours = {hours}")
print(f"Predicted Score = {own_pred[0]}")
```

No of Hours = 9.25

Predicted Score = 93.69173248737535

## Step 7: Evaluate the Model

It is important to evaluate the performance of algorithm to compare how well different algorithms perform on a particular dataset.

### 1. Mean Absolute Error

In [21]: `from sklearn import metrics`

```
print('Mean Absolute Error:', metrics.mean_absolute_error(y_test, y_pred))
```

Mean Absolute Error: 4.183859899002975

### 2. Max Error

In [22]: `print('Max Error:', metrics.max_error(y_test, y_pred))`

Max Error: 6.732260779489849

### 3. Mean Squared Error

In [23]: `print('Mean Squared Error:', metrics.mean_squared_error(y_test, y_pred))`

Mean Squared Error: 21.598769307217406

In [ ]: